

Assignment-15: Try various CNN networks on MNIST dataset [M]

Three different architecture of CNN network on MNIST datasets. MNIST datasets contains handwritten images .

Objective:

- 1) 3_ConvNets with kernel 3x3
- 2) 5_ConvNets with kernel 5x5
- 3) 7_ConvnNets with kernel 2x2

```
In [0]: # Credits: https://github.com/keras-team/keras/blob/master/examples/mnist_cnn.py

from datetime import datetime
from __future__ import print_function
import keras
from keras.datasets import mnist
from keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten
from keras.layers import Conv2D, MaxPooling2D
from keras import backend as K

from keras.initializers import he_normal
from keras.layers.normalization import BatchNormalization

batch_size = 128
num_classes = 10
epochs = 12

# input image dimensions
img_rows, img_cols = 28, 28
```

The data, split between train and test sets

```
In [0]: #
(x_train, y_train), (x_test, y_test) = mnist.load_data()
```

```
In [12]: if K.image_data_format() == 'channels_first':
    x_train = x_train.reshape(x_train.shape[0], 1, img_rows, img_cols)
    x_test = x_test.reshape(x_test.shape[0], 1, img_rows, img_cols)
    input_shape = (1, img_rows, img_cols)
else:
    x_train = x_train.reshape(x_train.shape[0], img_rows, img_cols, 1)
    x_test = x_test.reshape(x_test.shape[0], img_rows, img_cols, 1)
    input_shape = (img_rows, img_cols, 1)

x_train = x_train.astype('float32')
x_test = x_test.astype('float32')
x_train /= 255
x_test /= 255
print('x_train shape:', x_train.shape)
print(x_train.shape[0], 'train samples')
print(x_test.shape[0], 'test samples')

# convert class vectors to binary class matrices
y_train = keras.utils.to_categorical(y_train, num_classes)
y_test = keras.utils.to_categorical(y_test, num_classes)
```

```
x_train shape: (60000, 28, 28, 1)
60000 train samples
10000 test samples
```

```
In [13]: print(y_train.shape)
```

```
(60000, 10)
```

```
In [0]: %matplotlib inline
import matplotlib.pyplot as plt
import numpy as np
import time
# https://gist.github.com/greydanus/f6eee59eaf1d90fcb3b534a25362cea4
# https://stackoverflow.com/a/14434334
# this function is used to update the plots for each epoch and error
def plt_dynamic(x, vy, ty):
    fig = plt.figure( facecolor='y', edgecolor='k')
    plt.plot(x, vy, 'b', label="Validation Loss")
    plt.plot(x, ty, 'r', label="Train Loss")
    plt.xlabel('Epochs')
    plt.ylabel('Categorical Crossentropy Loss')
    plt.legend()
    plt.grid()
    plt.show()
```

1 Model 1:CNN with 3 ConvNet & 3x3 kernel size

```
In [17]: convnet3=Sequential() # Initializing the model

# First ConvNet
convnet3.add(Conv2D(32,kernel_size=(3,3),
                    activation='relu',
                    input_shape=input_shape))

convnet3.add(Conv2D(64,kernel_size=(3,3),
                    activation='relu'))

convnet3.add(Dropout(0.25))

convnet3.add(Conv2D(128,kernel_size=(3,3),
                    activation='relu'))
#maxpooling by (2,2) ,dropout,flattening
convnet3.add(MaxPooling2D(pool_size=(2,2)))
convnet3.add(Dropout(0.25))
convnet3.add(Flatten())

#hidden_layer
convnet3.add(Dense(256,
                  activation='relu',
                  kernel_initializer=he_normal(seed=None)))
convnet3.add(Dropout(0.5))
convnet3.add(Dense(num_classes,activation='softmax'))
print(convnet3.summary())
```

Layer (type)	Output Shape	Param #
conv2d_7 (Conv2D)	(None, 26, 26, 32)	320
conv2d_8 (Conv2D)	(None, 24, 24, 64)	18496
dropout_6 (Dropout)	(None, 24, 24, 64)	0
conv2d_9 (Conv2D)	(None, 22, 22, 128)	73856
max_pooling2d_3 (MaxPooling2D)	(None, 11, 11, 128)	0
dropout_7 (Dropout)	(None, 11, 11, 128)	0
flatten_3 (Flatten)	(None, 15488)	0
dense_3 (Dense)	(None, 256)	3965184
dropout_8 (Dropout)	(None, 256)	0
dense_4 (Dense)	(None, 10)	2570
Total params: 4,060,426		
Trainable params: 4,060,426		
Non-trainable params: 0		
None		

```
In [24]: #Model compilation
convnet3.compile(optimizer=keras.optimizers.Adam(),
                 loss=keras.losses.categorical_crossentropy,
                 metrics=['accuracy'])
convnet3_history=convnet3.fit(x_train,y_train,batch_size=batch_size,
                             epochs=epochs,
                             verbose=1,
                             validation_data=(x_test, y_test))
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 555s 9ms/step - loss: 0.1380 -
acc: 0.9585 - val_loss: 0.0364 - val_acc: 0.9878

Epoch 2/12

60000/60000 [=====] - 555s 9ms/step - loss: 0.0576 -
acc: 0.9823 - val_loss: 0.0297 - val_acc: 0.9906

Epoch 3/12

60000/60000 [=====] - 555s 9ms/step - loss: 0.0420 -
acc: 0.9872 - val_loss: 0.0289 - val_acc: 0.9908

Epoch 4/12

60000/60000 [=====] - 553s 9ms/step - loss: 0.0333 -
acc: 0.9898 - val_loss: 0.0252 - val_acc: 0.9923

Epoch 5/12

60000/60000 [=====] - 557s 9ms/step - loss: 0.0290 -
acc: 0.9909 - val_loss: 0.0218 - val_acc: 0.9931

Epoch 6/12

60000/60000 [=====] - 555s 9ms/step - loss: 0.0247 -
acc: 0.9919 - val_loss: 0.0231 - val_acc: 0.9930

Epoch 7/12

60000/60000 [=====] - 557s 9ms/step - loss: 0.0219 -
acc: 0.9929 - val_loss: 0.0259 - val_acc: 0.9920

Epoch 8/12

60000/60000 [=====] - 555s 9ms/step - loss: 0.0197 -
acc: 0.9937 - val_loss: 0.0237 - val_acc: 0.9929

Epoch 9/12

60000/60000 [=====] - 563s 9ms/step - loss: 0.0172 -
acc: 0.9943 - val_loss: 0.0209 - val_acc: 0.9938

Epoch 10/12

60000/60000 [=====] - 561s 9ms/step - loss: 0.0142 -
acc: 0.9953 - val_loss: 0.0285 - val_acc: 0.9916

Epoch 11/12

60000/60000 [=====] - 564s 9ms/step - loss: 0.0149 -
acc: 0.9956 - val_loss: 0.0222 - val_acc: 0.9940

Epoch 12/12


60000/60000 [=====] - 562s 9ms/step - loss: 0.0131 -
acc: 0.9957 - val_loss: 0.0259 - val_acc: 0.9938

```
In [32]: #evaluating model

score=convnet3.evaluate(x_test,y_test,verbose=0)
test_score3=score[0]
test_accuracy3=score[1]
train_accuracy3=max(convnet3_history.history['acc'])
print('test score : ',test_score3)
print('test sccuracy : ',test_accuracy3)
# error plot
x=list(range(1,epochs+1))
vy=convnet3_history.history['val_loss'] #validation loss
ty=convnet3_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

test score : 0.025890892835492695

test sccuracy : 0.9938



2 Model2:CNN with 5 ConvNet & kernel_size=(5x5)

5 convNet followed by maxpooling(2,2) and dropout

```
In [39]: convnet5=Sequential() # Initializing the model

# First ConvNet
convnet5.add(Conv2D(32,kernel_size=(5,5),
                    activation='relu',
                    padding='same',
                    input_shape=input_shape))

convnet5.add(Conv2D(64,kernel_size=(5,5),
                    padding='same',
                    activation='relu'))#Second Convnet
convnet5.add(MaxPooling2D(pool_size=(2,2)))
convnet5.add(Dropout(0.25))

convnet5.add(Conv2D(96,kernel_size=(5,5),
                    padding='same',
                    activation='relu')) # 3rd ConvNet
#maxpooling by (2,2) ,dropout,flattening
convnet5.add(MaxPooling2D(pool_size=(2,2)))
convnet5.add(Dropout(0.25))

convnet5.add(Conv2D(128,kernel_size=(5,5),
                    padding='same',
                    activation='relu'))#fourth Convnet
convnet5.add(MaxPooling2D(pool_size=(2,2)))
convnet5.add(Dropout(0.25))
convnet5.add(Conv2D(164,kernel_size=(5,5),
                    padding='same',
                    activation='relu'))#fifth Convnet
convnet5.add(MaxPooling2D(pool_size=(2,2)))
convnet5.add(Dropout(0.25))
convnet5.add(Flatten())

#hidden_Layer
convnet5.add(Dense(256,
                  activation='relu',
                  kernel_initializer=he_normal(seed=None)))
convnet5.add(BatchNormalization())
convnet5.add(Dropout(0.5))
convnet5.add(Dense(num_classes,activation='softmax'))
print(convnet5.summary())
```

Layer (type)	Output Shape	Param #
conv2d_20 (Conv2D)	(None, 28, 28, 32)	832
conv2d_21 (Conv2D)	(None, 28, 28, 64)	51264
max_pooling2d_14 (MaxPooling)	(None, 14, 14, 64)	0
dropout_17 (Dropout)	(None, 14, 14, 64)	0
conv2d_22 (Conv2D)	(None, 14, 14, 96)	153696
max_pooling2d_15 (MaxPooling)	(None, 7, 7, 96)	0
dropout_18 (Dropout)	(None, 7, 7, 96)	0
conv2d_23 (Conv2D)	(None, 7, 7, 128)	307328
max_pooling2d_16 (MaxPooling)	(None, 3, 3, 128)	0
dropout_19 (Dropout)	(None, 3, 3, 128)	0
conv2d_24 (Conv2D)	(None, 3, 3, 164)	524964
max_pooling2d_17 (MaxPooling)	(None, 1, 1, 164)	0
dropout_20 (Dropout)	(None, 1, 1, 164)	0
flatten_4 (Flatten)	(None, 164)	0
dense_5 (Dense)	(None, 256)	42240
batch_normalization_1 (Batch Normalization)	(None, 256)	1024
dropout_21 (Dropout)	(None, 256)	0
dense_6 (Dense)	(None, 10)	2570
Total params: 1,083,918		
Trainable params: 1,083,406		
Non-trainable params: 512		
None		

```
In [42]: #Model compilation
start = datetime.now()
convnet5.compile(optimizer=keras.optimizers.Adam(),
                 loss=keras.losses.categorical_crossentropy,
                 metrics=['accuracy'])
convnet5_history=convnet5.fit(x_train,y_train,batch_size=batch_size,
                             epochs=epochs,
                             verbose=1,
                             validation_data=(x_test, y_test))
print("Time taken to run this cell :", datetime.now() - start)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 983s 16ms/step - loss: 0.6764
- acc: 0.7753 - val_loss: 0.2016 - val_acc: 0.9378

Epoch 2/12

60000/60000 [=====] - 981s 16ms/step - loss: 0.2034
- acc: 0.9396 - val_loss: 0.0988 - val_acc: 0.9679

Epoch 3/12

60000/60000 [=====] - 989s 16ms/step - loss: 0.1265
- acc: 0.9634 - val_loss: 0.0604 - val_acc: 0.9821

Epoch 4/12

60000/60000 [=====] - 985s 16ms/step - loss: 0.0895
- acc: 0.9740 - val_loss: 0.0415 - val_acc: 0.9863

Epoch 5/12

60000/60000 [=====] - 977s 16ms/step - loss: 0.0737
- acc: 0.9792 - val_loss: 0.0374 - val_acc: 0.9872

Epoch 6/12

60000/60000 [=====] - 970s 16ms/step - loss: 0.0607
- acc: 0.9823 - val_loss: 0.0290 - val_acc: 0.9909

Epoch 7/12

60000/60000 [=====] - 970s 16ms/step - loss: 0.0521
- acc: 0.9851 - val_loss: 0.0277 - val_acc: 0.9921

Epoch 8/12

60000/60000 [=====] - 972s 16ms/step - loss: 0.0494
- acc: 0.9860 - val_loss: 0.0279 - val_acc: 0.9917

Epoch 9/12

60000/60000 [=====] - 972s 16ms/step - loss: 0.0423
- acc: 0.9872 - val_loss: 0.0291 - val_acc: 0.9910

Epoch 10/12

60000/60000 [=====] - 985s 16ms/step - loss: 0.0390
- acc: 0.9885 - val_loss: 0.0289 - val_acc: 0.9921

Epoch 11/12

60000/60000 [=====] - 982s 16ms/step - loss: 0.0356
- acc: 0.9902 - val_loss: 0.0199 - val_acc: 0.9938

Epoch 12/12


60000/60000 [=====] - 981s 16ms/step - loss: 0.0320
- acc: 0.9905 - val_loss: 0.0234 - val_acc: 0.9930

Time taken to run this cell : 3:15:47.690485


```
In [43]: #evaluating model
score=convnet5.evaluate(x_test,y_test,verbose=0)
test_score5=score[0]
test_accuracy5=score[1]
train_accuracy5=max(convnet5_history.history['acc'])
print('test score : ',test_score5)
print('test Accuracy : ',test_accuracy5)
# error plot
x=list(range(1,epochs+1))
vy=convnet5_history.history['val_loss'] #validation loss
ty=convnet5_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

test score : 0.023384571030837107

test Accuracy : 0.993



3 Model3:CNN with 7 ConvNet & kernel_size=(2x2)

5 convNet followed by maxpooling(2,2) and dropout

```

In [49]: convnet7=Sequential() # Initializing the model

# First ConvNet
convnet7.add(Conv2D(16,kernel_size=(2,2),
                    activation='relu',
                    padding='same',strides=(1,1),
                    input_shape=input_shape))

convnet7.add(Conv2D(32,kernel_size=(2,2),
                    padding='same',strides=(2,2),
                    activation='relu'))#Second Convnet
#convnet7.add(MaxPooling2D(pool_size=(2,2)))
#convnet7.add(Dropout(0.25))

convnet7.add(Conv2D(64,kernel_size=(2,2),
                    padding='same',
                    activation='relu')) # 3rd ConvNet
#maxpooling by (2,2) ,dropout,flattening
#convnet7.add(MaxPooling2D(pool_size=(2,2)))
convnet7.add(Dropout(0.15))

convnet7.add(Conv2D(96,kernel_size=(2,2),
                    padding='same',
                    activation='relu'))#fourth Convnet
convnet7.add(MaxPooling2D(pool_size=(2,2)))
convnet7.add(Dropout(0.39))
convnet7.add(Conv2D(128,kernel_size=(2,2),
                    padding='same',
                    activation='relu'))#fifth Convnet
convnet7.add(MaxPooling2D(pool_size=(2,2)))
convnet7.add(Dropout(0.3))
convnet7.add(Conv2D(164,kernel_size=(2,2),
                    padding='same',
                    activation='relu'))#sixth Convnet
convnet7.add(Conv2D(164,kernel_size=(2,2),
                    padding='same',strides=(1,1),
                    activation='relu'))#seventh Convnet

convnet7.add(MaxPooling2D(pool_size=(2,2)))
convnet7.add(Dropout(0.4))
convnet7.add(Flatten())

#hidden_layer
convnet7.add(Dense(256,
                    activation='relu',
                    kernel_initializer=he_normal(seed=None)))#1 hidden layer
convnet7.add(BatchNormalization())
convnet7.add(Dropout(0.5))
convnet7.add(Dense(148,
                    activation='relu',
                    kernel_initializer=he_normal(seed=None)))#2 hidden layer
convnet7.add(BatchNormalization())
convnet7.add(Dropout(0.5))
convnet7.add(Dense(128,
                    activation='relu',

```

```
        kernel_initializer=he_normal(seed=None)))#3 hidden layer
convnet7.add(BatchNormalization())
convnet7.add(Dropout(0.5))
convnet7.add(Dense(num_classes,activation='softmax'))
print(convnet7.summary())
```

Layer (type)	Output Shape	Param #
conv2d_56 (Conv2D)	(None, 28, 28, 16)	80
conv2d_57 (Conv2D)	(None, 14, 14, 32)	2080
conv2d_58 (Conv2D)	(None, 14, 14, 64)	8256
dropout_41 (Dropout)	(None, 14, 14, 64)	0
conv2d_59 (Conv2D)	(None, 14, 14, 96)	24672
max_pooling2d_37 (MaxPooling)	(None, 7, 7, 96)	0
dropout_42 (Dropout)	(None, 7, 7, 96)	0
conv2d_60 (Conv2D)	(None, 7, 7, 128)	49280
max_pooling2d_38 (MaxPooling)	(None, 3, 3, 128)	0
dropout_43 (Dropout)	(None, 3, 3, 128)	0
conv2d_61 (Conv2D)	(None, 3, 3, 164)	84132
conv2d_62 (Conv2D)	(None, 3, 3, 164)	107748
max_pooling2d_39 (MaxPooling)	(None, 1, 1, 164)	0
dropout_44 (Dropout)	(None, 1, 1, 164)	0
flatten_6 (Flatten)	(None, 164)	0
dense_11 (Dense)	(None, 256)	42240
batch_normalization_5 (Batch Normalization)	(None, 256)	1024
dropout_45 (Dropout)	(None, 256)	0
dense_12 (Dense)	(None, 148)	38036
batch_normalization_6 (Batch Normalization)	(None, 148)	592
dropout_46 (Dropout)	(None, 148)	0
dense_13 (Dense)	(None, 128)	19072
batch_normalization_7 (Batch Normalization)	(None, 128)	512
dropout_47 (Dropout)	(None, 128)	0
dense_14 (Dense)	(None, 10)	1290
Total params: 379,014		
Trainable params: 377,950		
Non-trainable params: 1,064		

None

```
In [50]: #Model compilation
start=datetime.now()
convnet7.compile(optimizer=keras.optimizers.Adam(),
                 loss=keras.losses.categorical_crossentropy,
                 metrics=['accuracy'])
convnet7_history=convnet7.fit(x_train,y_train,batch_size=batch_size,
                             epochs=epochs,
                             verbose=1,
                             validation_data=(x_test, y_test))
print("Time taken to run this cell :", datetime.now() - start)
```

Train on 60000 samples, validate on 10000 samples

Epoch 1/12

60000/60000 [=====] - 168s 3ms/step - loss: 1.9486 -
acc: 0.3600 - val_loss: 0.8981 - val_acc: 0.6529

Epoch 2/12

60000/60000 [=====] - 162s 3ms/step - loss: 0.4479 -
acc: 0.8694 - val_loss: 0.0940 - val_acc: 0.9737

Epoch 3/12

60000/60000 [=====] - 163s 3ms/step - loss: 0.1987 -
acc: 0.9485 - val_loss: 0.0649 - val_acc: 0.9835

Epoch 4/12

60000/60000 [=====] - 164s 3ms/step - loss: 0.1433 -
acc: 0.9640 - val_loss: 0.0442 - val_acc: 0.9873

Epoch 5/12

60000/60000 [=====] - 163s 3ms/step - loss: 0.1137 -
acc: 0.9715 - val_loss: 0.0675 - val_acc: 0.9833

Epoch 6/12

60000/60000 [=====] - 163s 3ms/step - loss: 0.1026 -
acc: 0.9748 - val_loss: 0.0373 - val_acc: 0.9905

Epoch 7/12

60000/60000 [=====] - 164s 3ms/step - loss: 0.0904 -
acc: 0.9783 - val_loss: 0.0384 - val_acc: 0.9900

Epoch 8/12

60000/60000 [=====] - 165s 3ms/step - loss: 0.0869 -
acc: 0.9790 - val_loss: 0.0333 - val_acc: 0.9918

Epoch 9/12

60000/60000 [=====] - 165s 3ms/step - loss: 0.0747 -
acc: 0.9818 - val_loss: 0.0298 - val_acc: 0.9928

Epoch 10/12

60000/60000 [=====] - 162s 3ms/step - loss: 0.0719 -
acc: 0.9820 - val_loss: 0.0302 - val_acc: 0.9931

Epoch 11/12

60000/60000 [=====] - 163s 3ms/step - loss: 0.0666 -
acc: 0.9830 - val_loss: 0.0275 - val_acc: 0.9932

Epoch 12/12

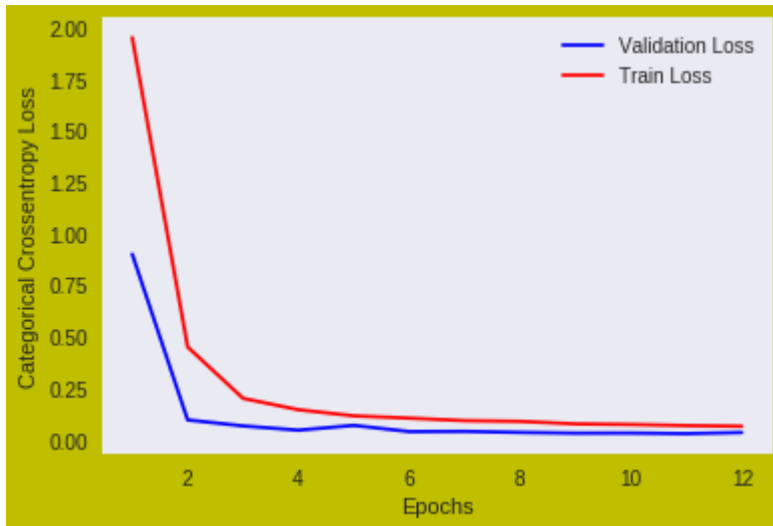
60000/60000 [=====] - 163s 3ms/step - loss: 0.0633 -
acc: 0.9844 - val_loss: 0.0330 - val_acc: 0.9926

Time taken to run this cell : 0:32:47.172971

```
In [51]: #evaluating model
score=convnet7.evaluate(x_test,y_test,verbose=0)
test_score7=score[0]
test_accuracy7=score[1]
train_accuracy7=max(convnet7_history.history['acc'])
print('test score : ',test_score7)
print('test Accuracy : ',test_accuracy7)
# error plot
x=list(range(1,epochs+1))
vy=convnet7_history.history['val_loss'] #validation loss
ty=convnet7_history.history['loss'] # train loss
plt_dynamic(x, vy, ty)
```

test score : 0.032985285274824125

test Accuracy : 0.9926



Observation

```
In [53]: from prettytable import PrettyTable
models=['3ConvNet with kernel 3x3',
        '5ConvNet with kernel 5x5',
        '7ConvNet with kernel 2x2']
training_accuracy=[train_accuracy3,train_accuracy5,train_accuracy7]
test_accuracy=[test_accuracy3,test_accuracy5,test_accuracy7]
INDEX = [1,2,3]
# Initializing prettytable
Model_Performance = PrettyTable()
# Adding columns
Model_Performance.add_column("INDEX.",INDEX)
Model_Performance.add_column("MODEL_NAME",models)
Model_Performance.add_column("TRAINING ACCURACY",training_accuracy)
Model_Performance.add_column("TESTING ACCURACY",test_accuracy)
#Model_Performance.add_column("TEST SCORE",test_score)

# Printing the Model_Performance
print(Model_Performance)
```

INDEX.	MODEL_NAME	TRAINING ACCURACY	TESTING ACCURACY
1	3ConvNet with kernel 3x3	0.9957166666666667	0.9938
2	5ConvNet with kernel 5x5	0.9905000000317892	0.993
3	7ConvNet with kernel 2x2	0.98441666666348775	0.9926

- from graphs in model1,model2 and model 3 model3 gives best plot amongst other model.
- Non_trainable params:

Model1=0

Model2=512

Model3=1064

- The Training time for model2 and model1 is very high comapratively model3.
- Model performance with three different CNN model is shown in above table.